



Neo4j Graph DB

{A step ahead towards non-RDBMS}

Presented By:
Mangal Dev



Agenda

- NOSQL
- Graph Database
Brew install neo4j
- Neo4J
- Samples





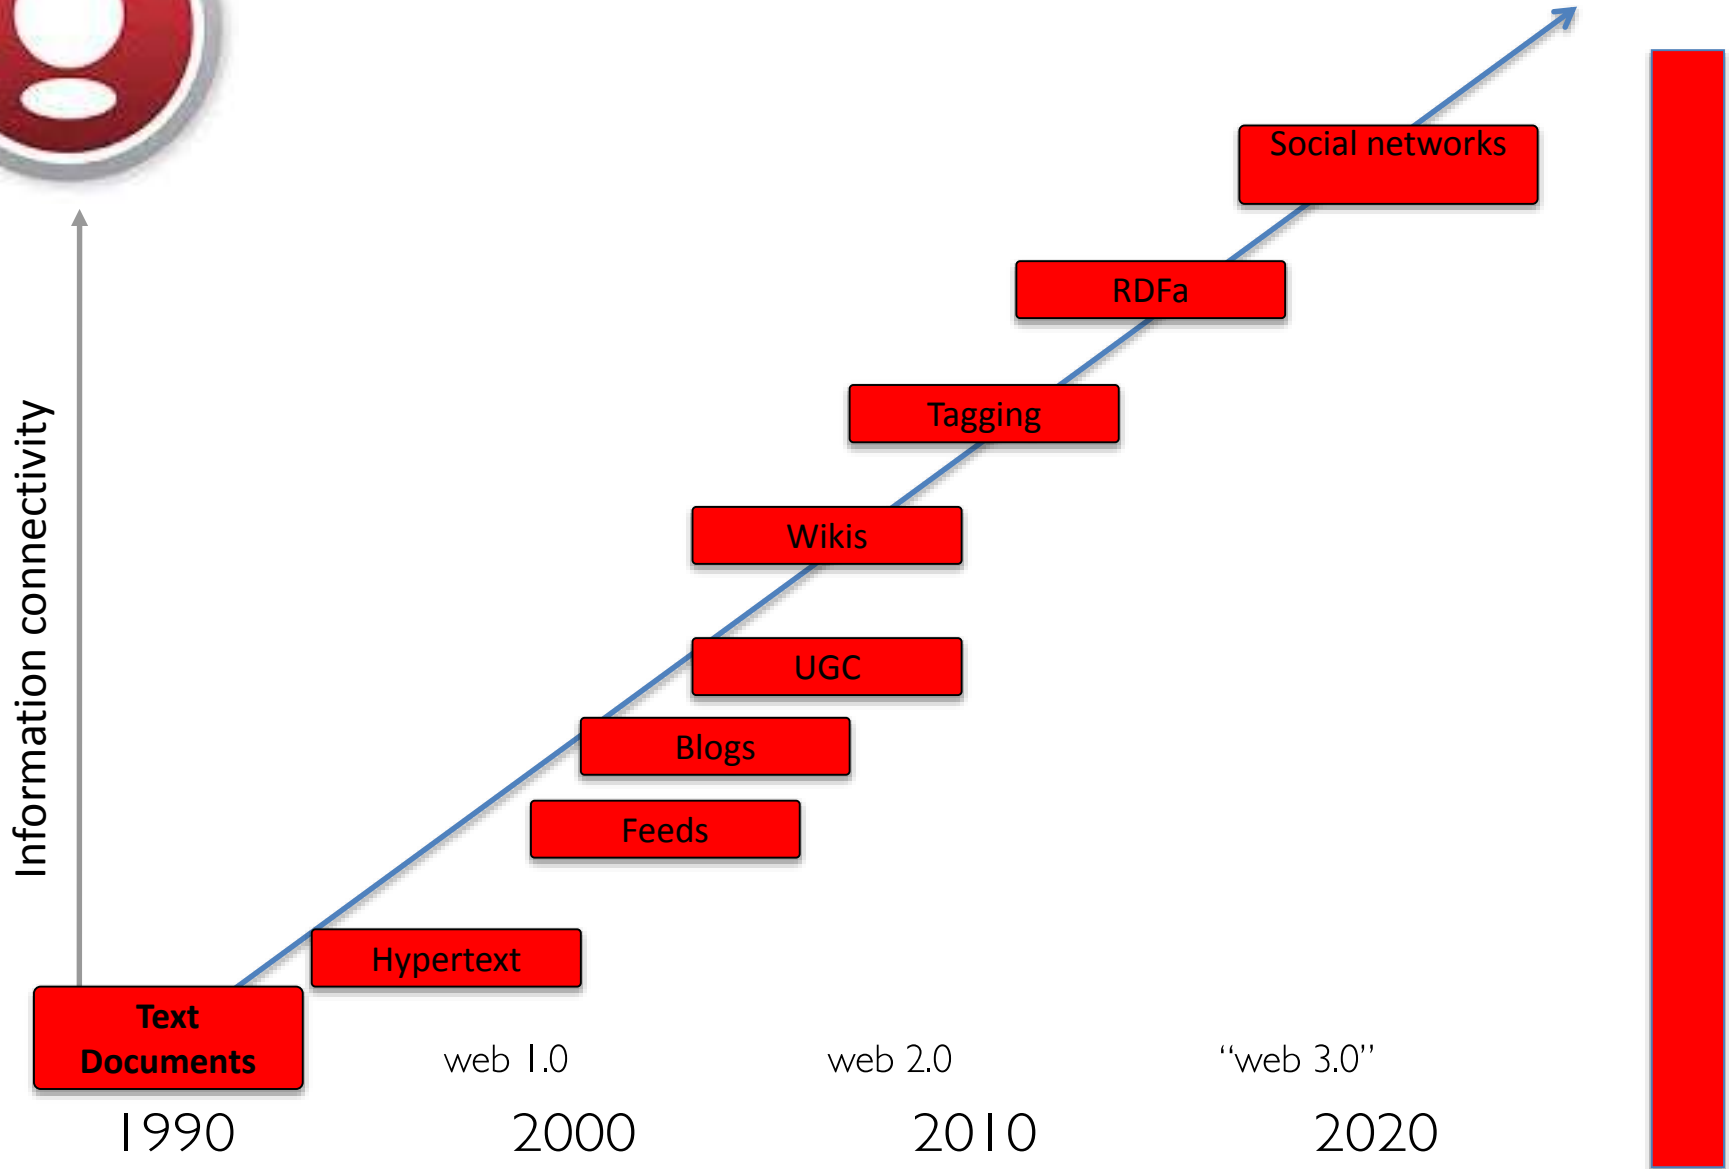
Not Only SQL

Why ?





Interdependency and complexity of data.



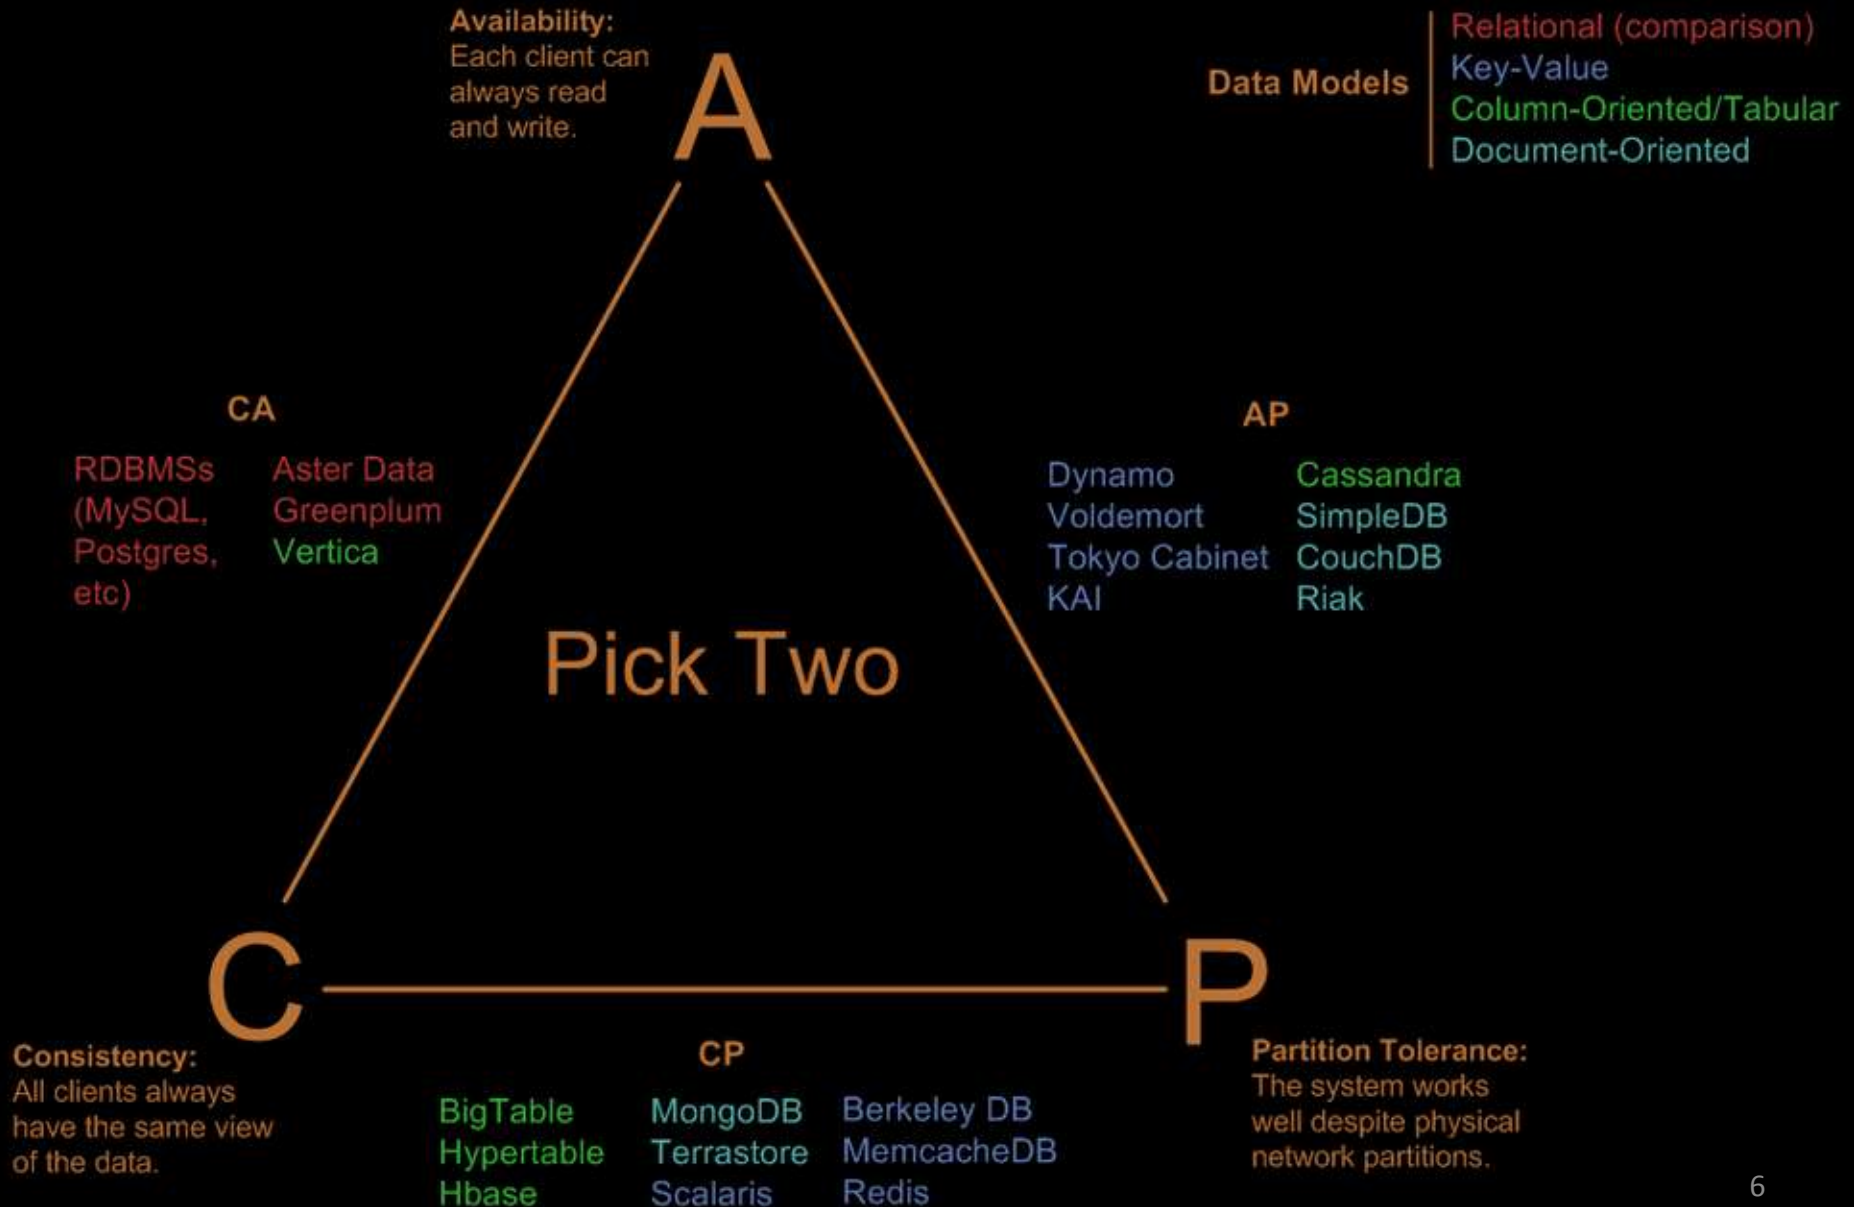


NOSQL Categories

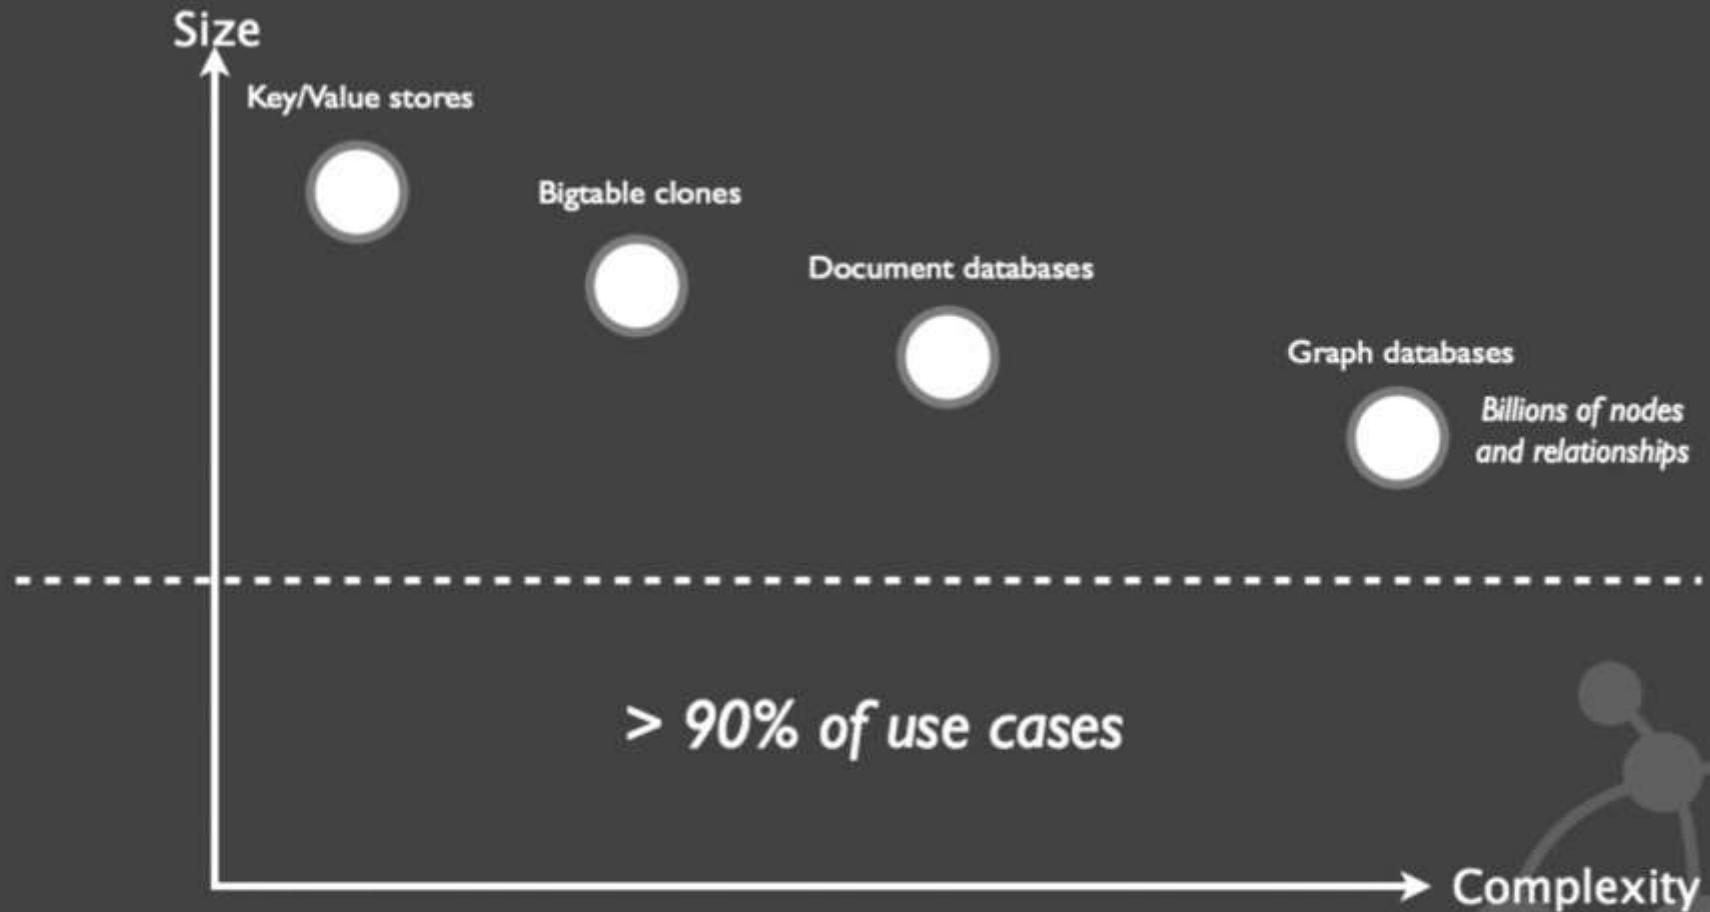
- **Key-Values Stores**
 - Dynamo
 - Riak
- **Column Family Stores**
 - Big Table -> Google
 - Cassandra -> Facebook
 - HBase -> Apache
- **Document Based DB**
 - CouchDB
 - MongoDB
- **Graph Databases.**
 - AllegroDB
 - FlockDB -> Twitter
 - InfiniteGraph
 - Neo4j
 - Sones



Visual Guide to NoSQL Systems



Scaling to size vs. Scaling to complexity





Why GraphDB ?

- Flexibility
 - No Normalization/De-normalization
 - Relational databases deal poorly with relationships.
- Agility
 - Schema free nature
 - Foreign Keys turns into dangling references
- Reciprocal Queries
- Usecases
 - SocialNetworks
 - BioInformatics
 - DataManagement
 - NetworkManagement
 - CloudManagement
 - GeoData

What does 'R' stand for in RDBMS ?



Graph Databases



sones



InfiniteGraph
Powered by Objectivity

**dex*



Neo4j



OrientDB®

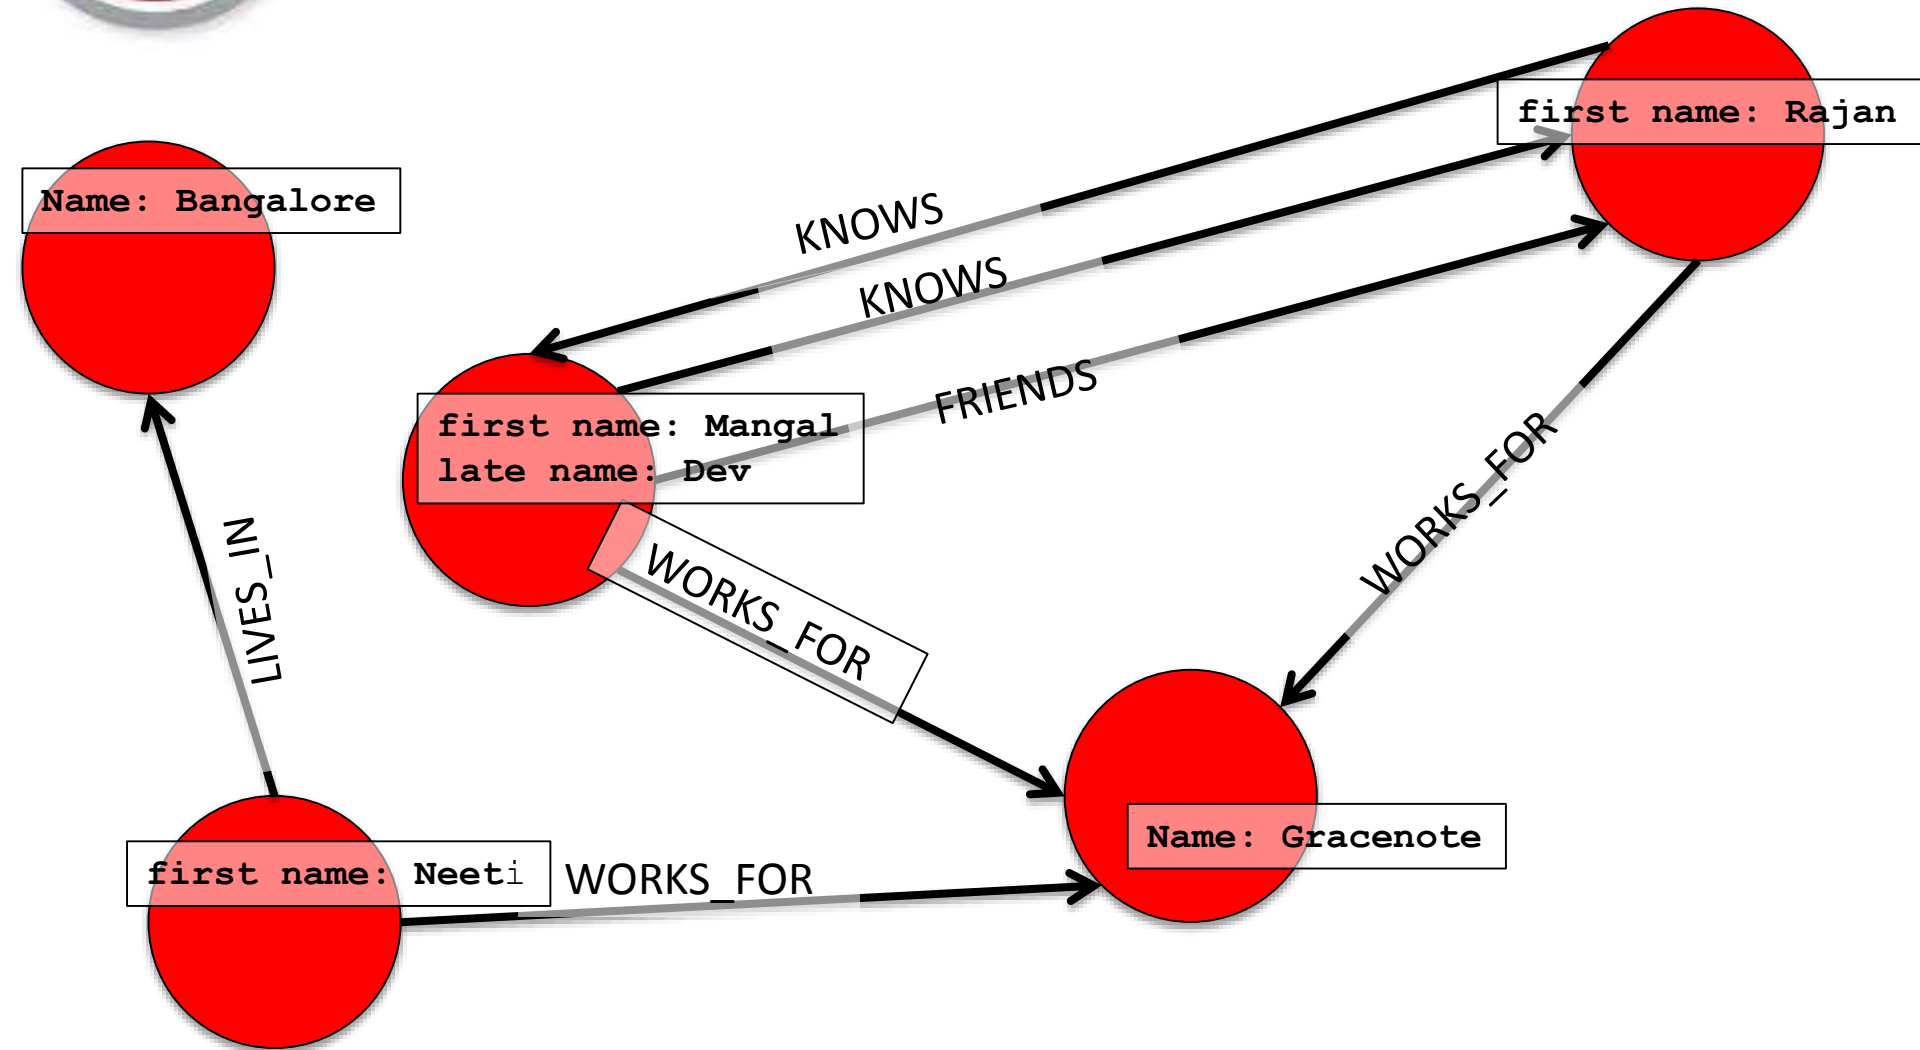


AllegroGraph

HYPERGRAPHDB






Property Graph Model





You know relational
now consider relationships...





Database full of linked nodes
Stores data as nodes and relationships



Some More Info..

- Widely Used
- Full ACID transactions
- Schema free, bottom-up data model design
- It's a stable and in active development since mid 2000
- High Performance, can cover upto 2 billion nodes.
- Lot of Bindings



How do you query and work with this graph database?

- Run it as
 - Embedded
 - Standalone
- Query it with
 - Cyhper from Neo4j Browser
 - Cypher from Java
 - Rest API
 - Rest API using Cypher query
 - Programmatically
 - Gremlin

Bindings



REST://





Cypher Query

- MATCH
 - Matches the graph pattern in the real graph.
- WHERE
 - Filters using predicates or anchors pattern elements.
- RETURN
 - Returns and projects result data, also handles aggregation.
- ORDER BY
 - Sorts the query result.
- SKIP/LIMIT
 - Paginates the query result.



Go to Console...

- Neo4j start
- Go to url :
 - <http://localhost:7474/browser/>



Cypher Basics



```
MATCH (a)-->(b)  
RETURN a, b;
```



Naming a Relationship



```
MATCH (a)-[r]->( )  
RETURN a.name, type (r);
```



Matching by Relationship



```
MATCH(a)-[:ACTED_IN]->(m)  
RETURN a.name, m.title;
```




Match Node Label

```
MATCH (tom:Person)
WHERE tom.name="Tom Hanks"
RETURN tom;
```

:Person – matches only nodes labeled
as Person



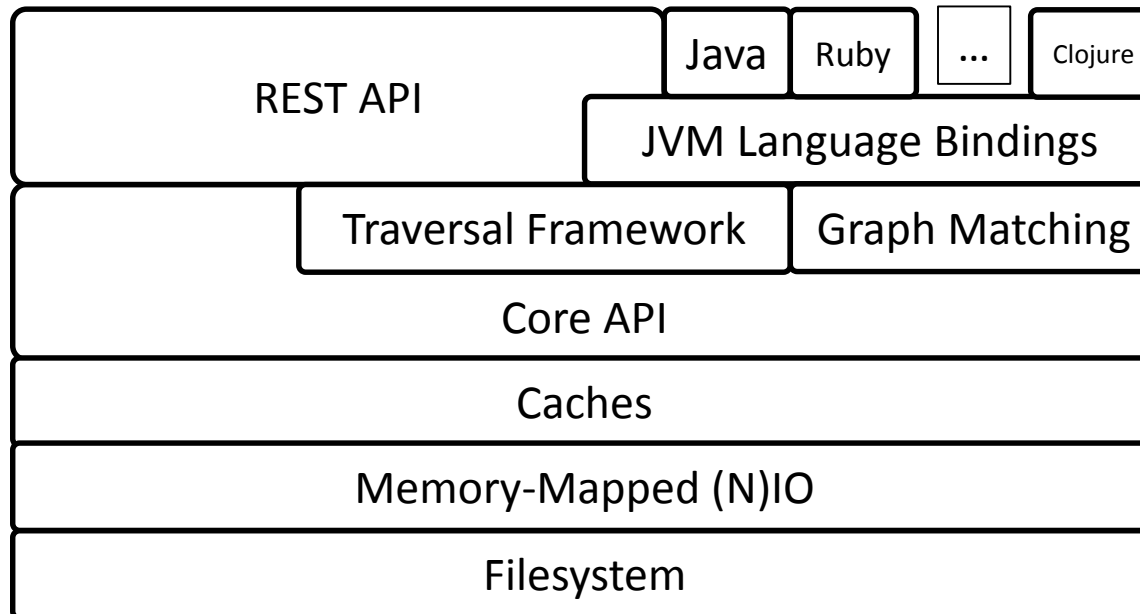
Variable length paths



$(a) - [*1..3] -> (b)$



Neo4j Logical Architecture





Graph Algorithms

- algorithm can have one of these values:
 - shortestPath
 - allSimplePaths
 - allPaths
 - dijkstra (optionally with cost_property and default_cost parameters)



Graph Algorithms

- `shortestPath()`

MATCH p =

ShortestPath

((keanu:Person)-[:KNOWS*]-(kevin:Person))

WHERE

keanu.name="Keanu Reeves"

and

kevin.name = "Kevin Bacon"

RETURN

length(p)



Indexing a Graph?

- Graphs are their own indexes!
- But sometimes we want short-cuts to well-known nodes
- Can do this in our own code
 - Just keep a reference to any interesting nodes

e.g.

```
CREATE INDEX ON :MOVIE(title);
```

Don't index every node!



Inbuilt Lucene Support

- Default index implementation for Neo4j
- Each index supports nodes *or* relationships
- Supports exact and regex-based matching
 - `Query('index_name','pattern')`
- Supports scoring
 - Number of hits in the index for a given item
 - Great for recommendations!



Scalability

- The HA component supports master-slave replication
- But Scaling graphs are **HARD.**
- Acc to Emil -> Superb performance till some billions of nodes



PERFORMANCE ANALYSIS

Table 2-1. Finding extended friends in a relational database versus efficient finding in Neo4j

Depth	RDBMS execution time (s)	Neo4j execution time (s)	Records returned
2	0.016	0.01	~2500
3	30.267	0.168	~110,000
4	1543.505	1.359	~600,000
5	Unfinished	2.132	~800,000



Pros and Cons

- Strengths
 - Powerful data model
 - Fast
 - For connected data, can be many orders of magnitude faster than RDBMS
- Weaknesses:
 - Sharding
 - Though they *can* scale reasonably well
 - And for some domains you can shard too!



ENOUGH THEORY!!!



Sample Exercises

- Find all Persons who killed other in Mahabharat.
- Find movies in which actor has acted along with directed it.
- the five actors who have acted in the most movies
- Shortest path between Kunti and Sehdev
- Find persons got killed by Kunti's sons



References

- http://www.neo4j.org/learn/online_course
- <http://www.infoq.com/articles/graph-nosql-neo4j>
- http://www.manning.com/partner/Neo4J_me_ap_ch01.pdf